

Qwen3.6-27B多模态推理

1. 基础环境

模型：

Hygon版本和沐曦版本

沐曦版本的量化模型（<https://www.modelscope.cn/models/metax-tech/Qwen3.6-27B-W8A8/files>）在容器qwen36_27b_zyf 容器内路径是/zhangyifan/qwen36-27 /workspace/model这里面是海光版本的模型

镜像：

#42.228.13.241:5000/jenkins/model_test_env/vllm:0.21.0-ubuntu22.04-dtk2604-py3.10-20260714-0235

没跑起来 镜像有问题

镜像

代码块

```
1 harbor.sourcefind.cn:5443/dcu/admin/base/vllm:0.18.1-ubuntu22.04-dtk26.04-py3.10
```

启动服务命令

代码块

```
1 export HIP_VISIBLE_DEVICES=4
2 export VLLM_HCU_USE_FLASH_ATTN=1
3 export VLLM_HCU_USE_CUSTOM_TOPK_TOPP_SAMPLER=1
4
5 vllm serve /workspace/model \
6     -tp 1 \
7     --port 8010 \
8     --trust-remote-code \
9     --max-num-batched-tokens 10240 \
10    -q slimquant_marlin \
11    --speculative-config.method mtp \
12    --speculative-config.num_speculative_tokens 3 \
```

```
13 --speculative-config.quantization "slimquant_marlin"
```

压测脚本



 vlm_speed_benchmark.py

压测命令

代码块

```
1 python3 vlm_speed_benchmark.py \  
2 --input /zhangyifan/qwen36-27/portable_2k_v1/video8.jsonl \  
3 --output /zhangyifan/qwen36-27/benchmark_video8.jsonl \  
4 --endpoints http://localhost:8010/v1/chat/completions \  
5 --model /workspace/model \  
6 --tokenizer-path /workspace/model \  
7 --workers 8 \  

```

2. 测试结果

启动客户端后服务端报错 修改压测脚本：

```

1 代码块 python3 - <<'EOF'
2  import pathlib
3  p = pathlib.Path("/zhangyifan/qwen36-27/vlm_speed_benchmark.py")
4  s = p.read_text()
5  old = '"do_sample_frames": True'
6  new = '"do_sample_frames": False'
7  if old in s:
8      p.write_text(s.replace(old, new))
9      print("patched:", p)
10 else:
11     print("pattern not found")
12 EOF
13 patched: /zhangyifan/qwen36-27/vlm_speed_benchmark.py

```

目的:

vLLM 0.18.1 + transformers 5.5.0 的双重采样越界: vLLM 先解码视频 → 已剪成 32 帧张量

transformers 又拿到 do_sample_frames=True → 按 metadata.total_num_frames=98 去对 32 帧二次采样

→ indices 顶到 37 > 32 → IndexError → HTTP 400

do_sample_frames=False 后, transformers 不进二次采样分支, vLLM 解出的帧原样传给模型, 不越界

压测结果:

```

root@localhost:/zhangyifan/qwen36-27# python3 vlm_speed_benchmark.py \
--input /zhangyifan/qwen36-27/portable_2k_v1/video_only.jsonl \
--output /zhangyifan/qwen36-27/benchmark_video2.jsonl \
--endpoints http://localhost:8010/v1/chat/completions \
--model /workspace/model \
--tokenizer-path /workspace/model \
--workers 8 \
--retries 2
{
  "portable_version": "hcc_vlm_speed_portable_v1",
  "authoritative_runner_sha256": "91584e23889abc31c760e72ac0b5e381180db1939bc2bc57817b814a8fe06f6b",
  "records": 501,
  "pass_records": 331,
  "nonpass_records": 170,
  "status_counts": {
    "excluded_complete_input_over_8192_tokens": 170,
    "pass": 331
  },
  "duration_sec": 1086.751,
  "records_per_sec": 0.304577,
  "prompt_tokens": 2033007,
  "completion_tokens": 48616,
  "total_tokens": 2081623,
  "prompt_tokens_per_sec": 1870.72,
  "completion_tokens_per_sec": 44.735,
  "total_tokens_per_sec": 1915.455,
  "latency_ms": {
    "min": 5451.728,
    "p50": 23524.117,
    "p90": 29143.4,
    "p95": 32823.711,
    "p99": 37155.956,
    "max": 42698.076
  },
  "input": "/zhangyifan/qwen36-27/portable_2k_v1/video_only.jsonl",
  "input_sha256": "4a10daf12d889876cc6d4adadd0e7f3440b5ff30cc1451587372ab5f2516500e",
  "output": "/zhangyifan/qwen36-27/benchmark_video2.jsonl",
  "workers": 8,
  "timeout": 480,
  "retries": 2,
  "repair_rounds": 0,
  "repair_retries": 0,
  "model": "/workspace/model",
  "endpoints": [
    "http://localhost:8010/v1/chat/completions"
  ],
  "per_endpoint": {
    "http://localhost:8010/v1/chat/completions": {
      "records": 501,
      "pass": 331,
      "prompt_tokens": 2033007,
      "completion_tokens": 48616,
      "total_tokens": 2081623,
      "latency_ms_p50": 23524.117,
      "latency_ms_p95": 32823.711
    }
  },
  "input_stable_sha256": "ae934a3edbf4ef8d2a1950b6ccbdbe22375c2857273bc4c208c0d126a3986398",
  "record_uid_sequence_sha256": "0a514dc6d8ab47407ee90e0c8b726534a9f98d63cd60c17c8d418aa2269c3c7f",
  "output_sha256": "0156695fa41f97cbff0c7ce760abf94a80ef01c79d55e192779e019a6f80fa4c"
}

```

$(501-170) / (18 \times 60) = 0.306 \text{ samples/s}$ 比客户测出来的0.078sam/s提升很多 当前对A800性能比为66% 原因分析：客户选择的量化版本不对 原始模型是fp8的 客户用nv的卡做的量化 更换为沐曦的量化版本后性能正常 模型链接已同步客户

单位：samples/s

括号里为对A800的16样本性能比

	0.18	0.21	A800
16样本性能	0.6837(148.5%)	1.549(335.3%)	0.4602

全量性能	0.6141 (133.4%)	1.3605(295.6%)	/
------	-----------------	----------------	---

注：服务端未开启缓存命中

```
{
  "portable_version": "hcc_vlm_speed_portable_v1",
  "authoritative_runner_sha256": "91584e23889abc31c760e72ac0b5e381180db1939bc2bc57817b814a8fe06f6",
  "records": 16,
  "pass_records": 16,
  "nonpass_records": 0,
  "status_counts": {
    "pass": 16
  },
  "duration_sec": 23.401,
  "records_per_sec": 0.683719,
  "prompt_tokens": 39545,
  "completion_tokens": 2264,
  "total_tokens": 41809,
  "prompt_tokens_per_sec": 1689.854,
  "completion_tokens_per_sec": 96.746,
  "total_tokens_per_sec": 1786.6,
  "latency_ms": {
    "min": 4010.489,
    "p50": 12560.615,
    "p90": 18088.114,
    "p95": 18088.114,
    "p99": 18239.338,
    "max": 18239.338
  },
  "input": "/zhangyifan/qwen36-27/portable_2k_v1/mixed16.jsonl",
  "input_sha256": "f4e4e74f6791c6595c0dd1cf1e671b256077fb2dbaee2c3b97d3019cab14d08e",
  "output": "/zhangyifan/qwen36-27/benchmark_mixed16.jsonl",
  "workers": 8,
  "timeout": 480,
  "retries": 2,
  "repair_rounds": 0,
  "repair_retries": 0,
  "model": "/workspace/model",
  "endpoints": [
    "http://localhost:8010/v1/chat/completions"
  ],
  "per_endpoint": {
    "http://localhost:8010/v1/chat/completions": {
      "records": 16,
      "pass": 16,
      "prompt_tokens": 39545,
      "completion_tokens": 2264,
      "total_tokens": 41809,
      "latency_ms_p50": 12560.615,
```

0.18 16样本